





<https://github.com/zuow13176/Machine-Learning/tree/main/Supervised%20Learning/Classification%20on%20Colab%20using%20MNIST%20data%20set>


 03_classification.ipynb

File Edit View Insert Runtime Tools Help [Cannot save changes](#)


+ Code + Text  Copy to Drive

 Open in Colab  Open in Kaggle

<> ▾ Setup



First, let's import a few common modules, ensure Matplotlib plots figures inline and prepare a function to save the figures. We also check that Python 3.5 or later is installed (although Python 2.x may work, it is deprecated so we strongly recommend you use Python 3 instead), as well as Scikit-Learn ≥ 0.20 .



```
# Python ≥3.5 is required
import sys
assert sys.version_info >= (3, 5)

# Is this notebook running on Colab or Kaggle?
IS_COLAB = "google.colab" in sys.modules
IS_KAGGLE = "kaggle_secrets" in sys.modules

# Scikit-Learn ≥0.20 is required
import sklearn
assert sklearn.__version__ >= "0.20"

# Common imports
import numpy as np
import os

# to make this notebook's output stable across runs
np.random.seed(42)

# To plot pretty figures
%matplotlib inline
import matplotlib as mpl
import matplotlib.pyplot as plt
```

Executing (0s) Cell > save_fig() > savefig() > savefig() > print_figure() > print_png()



```
# To plot pretty figures
%matplotlib inline
import matplotlib as mpl
import matplotlib.pyplot as plt
mpl.rc('axes', labelsize=14)
mpl.rc('xtick', labelsize=12)
mpl.rc('ytick', labelsize=12)

# Where to save the figures
PROJECT_ROOT_DIR = "."
CHAPTER_ID = "classification"
IMAGES_PATH = os.path.join(PROJECT_ROOT_DIR, "images", CHAPTER_ID)
os.makedirs(IMAGES_PATH, exist_ok=True)

def save_fig(fig_id, tight_layout=True, fig_extension="png", resolution=300):
    path = os.path.join(IMAGES_PATH, fig_id + "." + fig_extension)
    print("Saving figure", fig_id)
    if tight_layout:
        plt.tight_layout()
    plt.savefig(path, format=fig_extension, dpi=resolution)
```

▼ MNIST

Warning: since Scikit-Learn 0.24, `fetch_openml()` returns a Pandas DataFrame by default. To avoid this and keep the same code as in t we use `as_frame=False`.



```
[2] from sklearn.datasets import fetch_openml
mnist = fetch_openml('mnist_784', version=1, as_frame=False)
mnist.keys()
```

```
dict_keys(['data', 'target', 'frame', 'feature_names', 'target_names', 'DESCR', 'details', 'categories', 'url'])
```

```
[3] X, y = mnist["data"], mnist["target"]  
X.shape
```

```
(70000, 784)
```

```
[4] y.shape
```

```
(70000,)
```

```
▶ 28 * 28
```

```
👤 784
```

```
[6] %matplotlib inline  
import matplotlib as mpl  
import matplotlib.pyplot as plt  
  
some_digit = X[0]  
some_digit_image = some_digit.reshape(28, 28)  
plt.imshow(some_digit_image, cmap=mpl.cm.binary)  
plt.axis("off")  
  
save_fig("some_digit_plot")  
plt.show()
```

```
Saving figure some_digit_plot
```




'5'


```
[8] y = y.astype(np.uint8)
```

```
▶ def plot_digit(data):  
    image = data.reshape(28, 28)  
    plt.imshow(image, cmap = mpl.cm.binary,  
                interpolation="nearest")  
    plt.axis("off")
```

```
[10] # EXTRA  
def plot_digits(instances, images_per_row=10, **options):  
    size = 28  
    images_per_row = min(len(instances), images_per_row)  
    images = [instance.reshape(size,size) for instance in instances]  
    n_rows = (len(instances) - 1) // images_per_row + 1  
    row_images = []  
    n_empty = n_rows * images_per_row - len(instances)  
    images.append(np.zeros((size, size * n_empty)))  
    for row in range(n_rows):  
        rimages = images[row * images_per_row : (row + 1) * images_per_row]  
        row_images.append(np.concatenate(rimages, axis=1))  
    image = np.concatenate(row_images, axis=0)  
    plt.imshow(image, cmap = mpl.cm.binary, **options)  
    plt.axis("off")
```

```
plt.figure(figsize=(9,9))
example_images = X[:100]
plot_digits(example_images, images_per_row=10)
save_fig("more_digits_plot")
plt.show()
```

 Saving figure more_digits_plot




```
In [ ]: def plot_precision_recall_vs_threshold(precisions, recalls, thresholds):
plt.plot(thresholds, precisions[:-1], "b--", label="Precision", linewidth=2)
plt.plot(thresholds, recalls[:-1], "g-", label="Recall", linewidth=2)
plt.legend(loc="center right", fontsize=16) # Not shown in the book
plt.xlabel("Threshold", fontsize=16) # Not shown
plt.grid(True) # Not shown
plt.axis([-50000, 50000, 0, 1]) # Not shown

recall_90_precision = recalls[np.argmax(precisions >= 0.90)]
threshold_90_precision = thresholds[np.argmax(precisions >= 0.90)]

plt.figure(figsize=(8, 4))
plot_precision_recall_vs_threshold(precisions, recalls, thresholds) # Not shown
plt.plot([threshold_90_precision, threshold_90_precision], [0., 0.9], "r:") # Not shown
plt.plot([-50000, threshold_90_precision], [0.9, 0.9], "r:") # Not shown
plt.plot([-50000, threshold_90_precision], [recall_90_precision, recall_90_precision], "r:") # Not shown
plt.plot([threshold_90_precision, threshold_90_precision], [0.9], "ro") # Not shown
plt.plot([threshold_90_precision, threshold_90_precision], [recall_90_precision], "ro") # Not shown
save_fig("precision_recall_vs_threshold_plot") # Not shown
plt.show()
```


Saving figure precision_recall_vs_threshold_plot



```
In [ ]: recall_for_forest = tpr_forest[np.argmax(fpr_forest >= fpr_90)]

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, "b:", linewidth=2, label="SGD")
plot_roc_curve(fpr_forest, tpr_forest, "Random Forest")
plt.plot([fpr_90, fpr_90], [0., recall_90_precision], "r:")
plt.plot([0.0, fpr_90], [recall_90_precision, recall_90_precision], "r:")
plt.plot([fpr_90], [recall_90_precision], "ro")
plt.plot([fpr_90, fpr_90], [0., recall_for_forest], "r:")
plt.plot([fpr_90], [recall_for_forest], "ro")
plt.grid(True)
plt.legend(loc="lower right", fontsize=16)
save_fig("roc_curve_comparison_plot")
plt.show()
```


Saving figure roc_curve_comparison_plot



```
In [ ]: cl_a, cl_b = 3, 5
X_aa = X_train[(y_train == cl_a) & (y_train_pred == cl_a)]
X_ab = X_train[(y_train == cl_a) & (y_train_pred == cl_b)]
X_ba = X_train[(y_train == cl_b) & (y_train_pred == cl_a)]
X_bb = X_train[(y_train == cl_b) & (y_train_pred == cl_b)]

plt.figure(figsize=(8,8))
plt.subplot(221); plot_digits(X_aa[:25], images_per_row=5)
plt.subplot(222); plot_digits(X_ab[:25], images_per_row=5)
plt.subplot(223); plot_digits(X_ba[:25], images_per_row=5)
plt.subplot(224); plot_digits(X_bb[:25], images_per_row=5)
save_fig("error_analysis_digits_plot")
plt.show()
```

Saving figure error_analysis_digits_plot



```
In [ ]: from sklearn.linear_model import LogisticRegression
        from sklearn.model_selection import cross_val_score

        log_clf = LogisticRegression(solver="lbfgs", max_iter=1000, random_state=42)
        score = cross_val_score(log_clf, X_train_transformed, y_train, cv=3, verbose=3)
        score.mean()
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.1s remaining: 0.0s
```

```
[CV] .....
[CV] ..... , score=0.981, total= 0.1s
[CV] .....
[CV] ..... , score=0.985, total= 0.2s
[CV] .....
[CV] ..... , score=0.991, total= 0.2s
```

```
[Parallel(n_jobs=1)]: Done 2 out of 2 | elapsed: 0.3s remaining: 0.0s
[Parallel(n_jobs=1)]: Done 3 out of 3 | elapsed: 0.5s finished
```

```
Out[ ]: 0.9858333333333333
```

Over 98.5%, not bad for a first try! :) However, remember that we are using the "easy" dataset. You can try with the harder datasets would have to try multiple models, select the best ones and fine-tune them using cross-validation, and so on.

But you get the picture, so let's stop now, and just print out the precision/recall we get on the test set:

```
In [ ]: from sklearn.metrics import precision_score, recall_score

        X_test_transformed = preprocess_pipeline.transform(X_test)

        log_clf = LogisticRegression(solver="lbfgs", max_iter=1000, random_state=42)
        log_clf.fit(X_train_transformed, y_train)

        y_pred = log_clf.predict(X_test_transformed)

        print("Precision: {:.2f}%".format(100 * precision_score(y_test, y_pred)))
        print("Recall: {:.2f}%".format(100 * recall_score(y_test, y_pred)))

        Precision: 95.88%
        Recall: 97.89%
```

```
In [ ]:
```